

Travel Insurance Purchase Forecast

Zahra Fatah

2023-11-20

A tour & travels company is offering travel insurance package to their customers. The new insurance package also includes COVID cover. The company wants to know which customers would be interested to buy it based on their database history. The insurance was offered to some of the customers in 2019 and the given data has been extracted from the performance/sales of the package during that period. The data is provided for almost 2000 of its previous customers and the goal is to build a model that can predict if the customer will be interested to buy the travel insurance package.

```
#par( mfrow= c(3,2) )

draw_confusion_matrix <- function(cm) {

  total <- sum(cm$table)
  res <- as.numeric(cm$table)

  # Generate color gradients. Palettes come from RColorBrewer.
  greenPalette <- c("#F7FCF5", "#E5F5E0", "#C7E9C0", "#A1D99B", "#74C476", "#41AB5D", "#238B45", "#006D2C", "#003366")
  redPalette <- c("#FFF5F0", "#FEE0D2", "#FCBBA1", "#FC9272", "#FB6A4A", "#EF3B2C", "#CB181D", "#A50F15", "#670000")
  getColor <- function (greenOrRed = "green", amount = 0) {
    if (amount == 0)
      return("#FFFFFF")
    palette <- greenPalette
    if (greenOrRed == "red")
      palette <- redPalette
    colorRampPalette(palette)(100)[10 + ceiling(90 * amount / total)]
  }

  # set the basic layout
  layout(matrix(c(1,1,2)))
  par(mar=c(2,2,2,2))
  plot(c(100, 345), c(300, 450), type = "n", xlab="", ylab="", xaxt='n', yaxt='n')
  title('CONFUSION MATRIX', cex.main=2)

  # create the matrix
  classes = colnames(cm$table)
  rect(150, 430, 240, 370, col=getColor("green", res[1]))
  text(195, 435, classes[1], cex=1.2)
  rect(250, 430, 340, 370, col=getColor("red", res[3]))
  text(295, 435, classes[2], cex=1.2)
  text(125, 370, 'Predicted', cex=1.3, srt=90, font=2)
  text(245, 450, 'Actual', cex=1.3, font=2)
  rect(150, 305, 240, 365, col=getColor("red", res[2]))
  rect(250, 305, 340, 365, col=getColor("green", res[4]))
  text(140, 400, classes[1], cex=1.2, srt=90)
```

```

text(140, 335, classes[2], cex=1.2, srt=90)

# add in the cm results
text(195, 400, res[1], cex=1.6, font=2, col='white')
text(195, 335, res[2], cex=1.6, font=2, col='white')
text(295, 400, res[3], cex=1.6, font=2, col='white')
text(295, 335, res[4], cex=1.6, font=2, col='white')

# add in the specifics
plot(c(100, 0), c(100, 0), type = "n", xlab="", ylab="", main = "DETAILS", xaxt='n', yaxt='n')
text(10, 85, names(cm$byClass[1]), cex=1.2, font=2)
text(10, 70, round(as.numeric(cm$byClass[1]), 3), cex=1.2)
text(30, 85, names(cm$byClass[2]), cex=1.2, font=2)
text(30, 70, round(as.numeric(cm$byClass[2]), 3), cex=1.2)
text(50, 85, names(cm$byClass[5]), cex=1.2, font=2)
text(50, 70, round(as.numeric(cm$byClass[5]), 3), cex=1.2)
text(70, 85, names(cm$byClass[6]), cex=1.2, font=2)
text(70, 70, round(as.numeric(cm$byClass[6]), 3), cex=1.2)
text(90, 85, names(cm$byClass[7]), cex=1.2, font=2)
text(90, 70, round(as.numeric(cm$byClass[7]), 3), cex=1.2)

# add in the accuracy information
text(50, 35, names(cm$overall[1]), cex=1.5, font=2)
text(50, 20, round(as.numeric(cm$overall[1]), 3), cex=1.4)
#text(70, 35, names(cm$overall[2]), cex=1.5, font=2)
#text(70, 20, round(as.numeric(cm$overall[2]), 3), cex=1.4)
}

data=read.csv("TravelInsuranceData.csv",header=TRUE)
TravelInsuranceTest=read.csv("TravelInsuranceTest.csv",header=T)

Insurance_data= data[,-1] # Removing very first column as it was not necessary in the data analysis.
TravelInsuranceTest=TravelInsuranceTest[,-1]

Insurance_data$ChronicDiseases=as.factor(Insurance_data$ChronicDiseases)
Insurance_data$Employment.Type= as.factor(Insurance_data$Employment.Type)
Insurance_data$GraduateOrNot= as.factor(Insurance_data$GraduateOrNot)
Insurance_data$FrequentFlyer= as.factor(Insurance_data$FrequentFlyer)
Insurance_data$EverTravelledAbroad= as.factor(Insurance_data$EverTravelledAbroad)
Insurance_data$TravelInsurance= as.factor(Insurance_data$TravelInsurance)

TravelInsuranceTest$ChronicDiseases=as.factor(TravelInsuranceTest$ChronicDiseases)
TravelInsuranceTest$Employment.Type= as.factor(TravelInsuranceTest$Employment.Type)
TravelInsuranceTest$GraduateOrNot= as.factor(TravelInsuranceTest$GraduateOrNot)
TravelInsuranceTest$FrequentFlyer= as.factor(TravelInsuranceTest$FrequentFlyer)
TravelInsuranceTest$EverTravelledAbroad= as.factor(TravelInsuranceTest$EverTravelledAbroad)
TravelInsuranceTest$TravelInsurance= as.factor(TravelInsuranceTest$TravelInsurance)

```

1. Before we create a model, do some data cleaning, feature selection and exploratory data analysis.

```
unique(Insurance_data$ChronicDiseases)
```

```
## [1] 1 0
## Levels: 0 1
```

```

unique(Insurance_data$Employment.Type)

## [1] Government Sector      Private Sector/Self Employed
## Levels: Government Sector Private Sector/Self Employed

unique(Insurance_data$GraduateOrNot)

## [1] Yes No
## Levels: No Yes

unique(Insurance_data$FrequentFlyer)

## [1] No  Yes
## Levels: No Yes

unique(Insurance_data$EverTravelledAbroad)

## [1] No  Yes
## Levels: No Yes

unique(Insurance_data$TravelInsurance)

## [1] 0 1
## Levels: 0 1

unique(Insurance_data$AnnualIncome)

## [1] 400000 1250000 500000 700000 1150000 1300000 1350000 1450000 800000
## [10] 1400000 850000 1500000 1050000 350000 600000 900000 550000 300000
## [19] 750000 1100000 1200000 1000000 950000 1700000 1750000 650000 450000
## [28] 1800000 1550000 1650000

unique(Insurance_data$Age)

## [1] 31 34 28 25 33 26 32 29 35 30 27

unique(Insurance_data$FamilyMembers)

## [1] 6 7 4 3 8 9 5 2

skim_without_charts(Insurance_data)

```

Table 1: Data summary


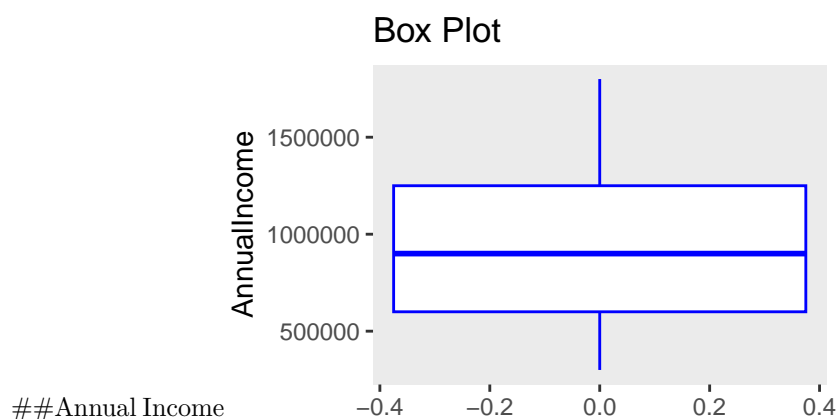
Name	Insurance_data
Number of rows	1887
Number of columns	9
Column type frequency:	
factor	6
numeric	3
Group variables	None

Variable type: factor

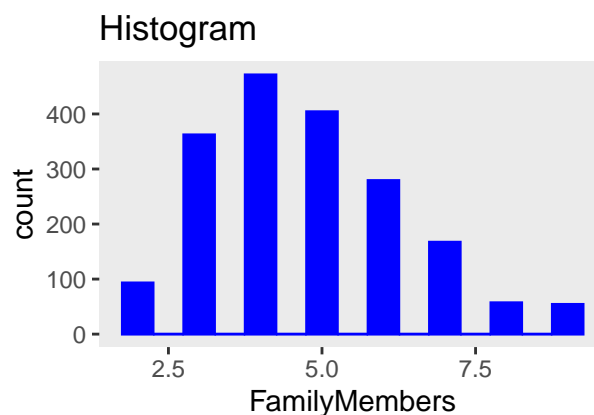

skim_variable	n_missing	complete_rate	ordered	n_unique	top_counts
Employment.Type	0	1	FALSE	2	Pri: 1352, Gov: 535
GraduateOrNot	0	1	FALSE	2	Yes: 1605, No: 282
ChronicDiseases	0	1	FALSE	2	0: 1359, 1: 528
FrequentFlyer	0	1	FALSE	2	No: 1495, Yes: 392
EverTravelledAbroad	0	1	FALSE	2	No: 1520, Yes: 367
TravelInsurance	0	1	FALSE	2	0: 1206, 1: 681

Variable type: numeric



skim_variable	n_missing	complete_rate	mean	sd	p0	p25	p50	p75	p100
Age	0	1	29.64	2.92	25	28	29	32	35
AnnualIncome	0	1	936062.53	376418.10	300000	600000	900000	1250000	1800000
FamilyMembers	0	1	4.75	1.62	2	4	5	6	9



Family members









Age



Categorical Variables

```
## [1] 0 1
```





```
## Levels: 0 1
```







Categorical Variables 2

```
## Warning: Unknown or uninitialised column: `(TravellInsurance)`.
```

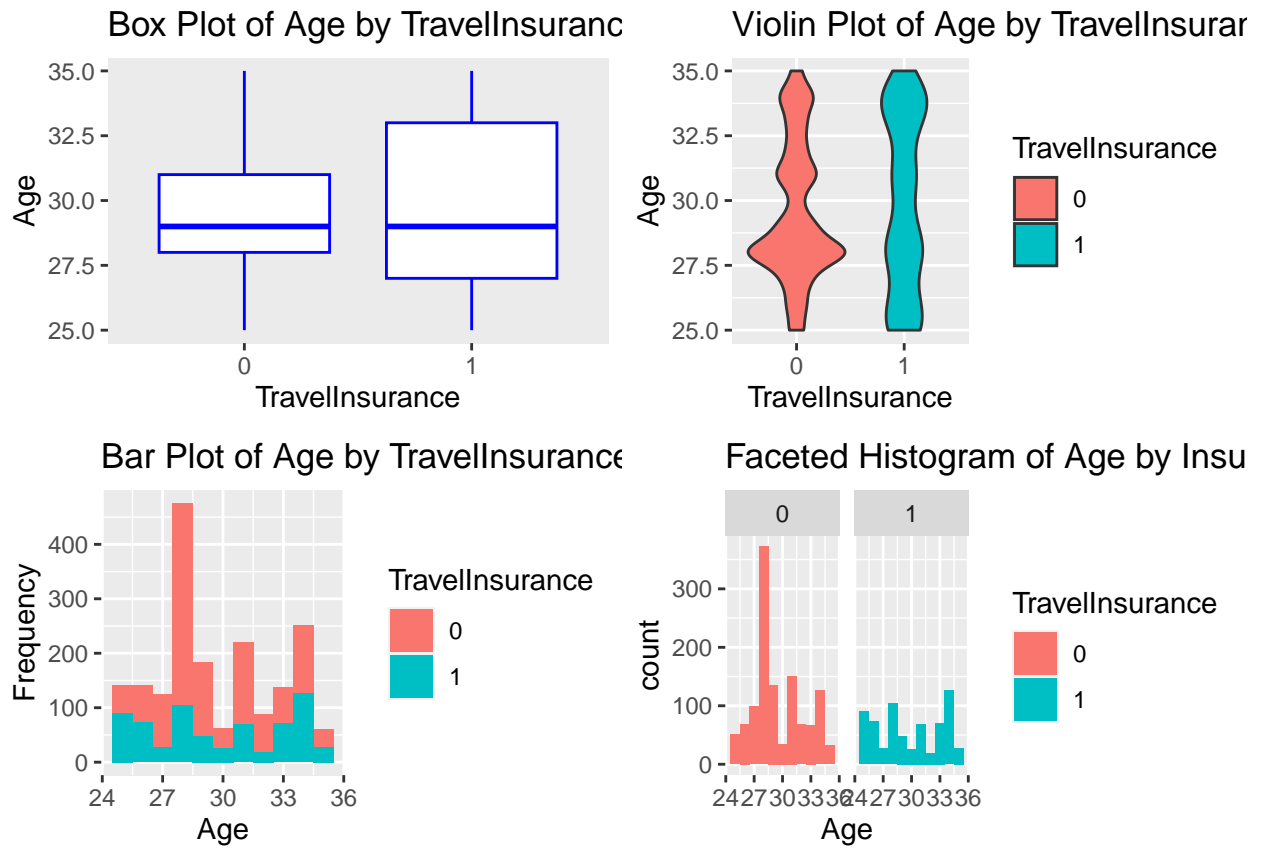
```
## NULL
```



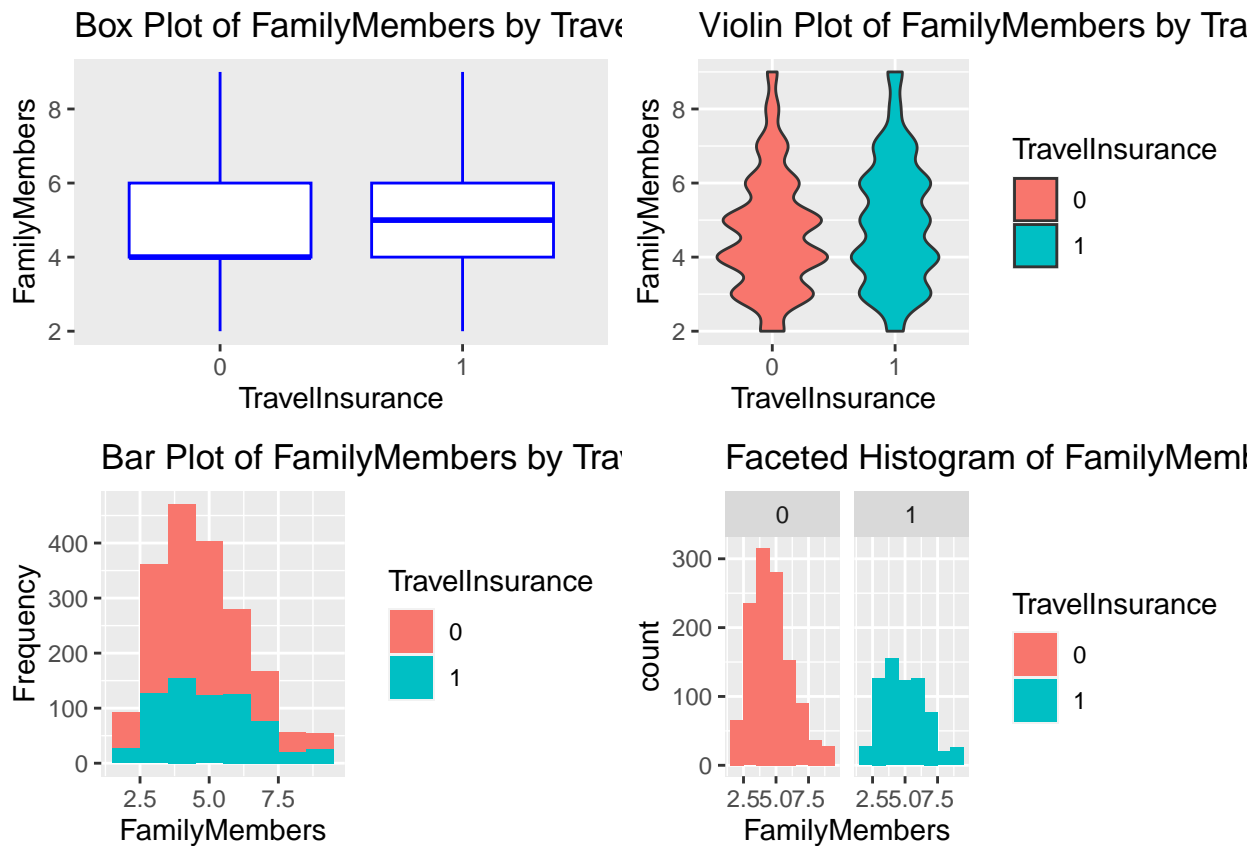
plot of AnnualIncome by Insurance



plot of Age by Insurance



plot of FamilyMembers by Insurance



```
#ggpairs(data = Insurance_data %>% select(TravelInsurance, Age, AnnualIncome, FamilyMembers))
```

2. Come up with a set of candidate methods that is suitable for the data.
3. Fit the models with training data.
4. Reduce the dimension of features by performing feature selection or dimension reduction.
5. Adjust the tuning parameters using cross-validation or model performance criteria such as error rate, AUC, etc.
6. Check the adequacy of the model fits and possibly revise the model.
7. Compare the models and choose your final model base on the prediction accuracy on the test data.

Logestic regression

2) Logistic Model

use all variables as predictor:

```
set.seed(100)

# Step 1: Split the data into training and testing sets
sample_index= sample(1:nrow(Insurance_data), 0.8 * nrow(Insurance_data))
train_data=Insurance_data[sample_index, ]
test_data=Insurance_data[-sample_index, ]
```

```

# Step 2: Train the logistic regression model
model= glm(TravelInsurance ~ Age + Employment.Type + GraduateOrNot + AnnualIncome + FamilyMembers +ChronicDiseases1 +FrequentFlyerYes +EverTravelledAbroad,
            family = binomial, data = train_data)
summary(model)

##
## Call:
## glm(formula = TravelInsurance ~ Age + Employment.Type + GraduateOrNot +
##      AnnualIncome + FamilyMembers + ChronicDiseases1 + FrequentFlyer +
##      EverTravelledAbroad, family = binomial, data = train_data)
##
## Coefficients:
##
##              Estimate Std. Error z value
## (Intercept)    -4.996e+00  7.137e-01  -7.001
## Age              6.275e-02  2.093e-02   2.997
## Employment.TypePrivate Sector/Self Employed  2.155e-01  1.493e-01   1.444
## GraduateOrNotYes -2.093e-01  1.739e-01  -1.203
## AnnualIncome     1.506e-06  1.965e-07   7.660
## FamilyMembers    1.403e-01  3.773e-02   3.719
## ChronicDiseases1  4.122e-02  1.366e-01   0.302
## FrequentFlyerYes  4.415e-01  1.547e-01   2.853
## EverTravelledAbroadYes  1.538e+00  1.720e-01   8.942
##
##              Pr(>|z|)
## (Intercept)    2.55e-12 ***
## Age            0.00272 **
## Employment.TypePrivate Sector/Self Employed  0.14880
## GraduateOrNotYes  0.22881
## AnnualIncome     1.85e-14 ***
## FamilyMembers    0.00020 ***
## ChronicDiseases1  0.76276
## FrequentFlyerYes  0.00433 **
## EverTravelledAbroadYes < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 1955.0  on 1508  degrees of freedom
## Residual deviance: 1603.1  on 1500  degrees of freedom
## AIC: 1621.1
##
## Number of Fisher Scoring iterations: 4

```

observation

-looking at the logistic model when trained on the training data, still statically significant variables and statically insignificant variables are the same.

Observation -Here Age, AnnualIncome, FamilyMembers, FrequentFlyerYes, EverTravelledAbroadYes are statically significant in determining weather the customer will purchase a travel insurance or not.

-likewise ChronicDiseases1, GraduateOrNotYes, Employment.TypePrivate Sector/Self Employed are not statically significant indicating that they are not important for customer in purchasing the travel insurance.

-for every one unit change in customers age, the log odd of purchasing travel insurance is increased by 7.29e-02

units

-for every one unit change in customers Annual income, the log odd of purchasing travel insurance is increased by 1.56e-06 units

-for every one unit change in customers number of family members in the family, the log odd of purchasing travel insurance is increased by 1.44e-01 units

```
set.seed(100)
# Step 3: Make predictions on the testing set
predictions_glm= predict(model, newdata = test_data, type = "response")
```

```
set.seed(100)

predicted_labels= ifelse(predictions_glm > 0.5, 1, 0)
```


let's see the confusion matrix

```
#####
pred=as.factor(predicted_labels)
cm_glm = confusionMatrix(pred,test_data$TravelInsurance,positive = "1")
cm_glm$positive
```

```
## [1] "1"
```

```
draw_confusion_matrix(cm_glm)
```

CONFUSION MATRIX



DETAILS

Sensitivity 0.546	Specificity 0.938	Precision 0.856	Recall 0.546	F1 0.667
Accuracy 0.78				

```
#####
table(predicted_labels,test_data$TravelInsurance)
```

```
##
```

```
## predicted_labels  0  1
```

```
##           0 212 69
##           1  14 83
```

```
#(accuracy_glm=(cm_glm[1,1]+cm_glm[2,2])/(cm_glm[1,1]+cm_glm[1,2]+cm_glm[2,1]+cm_glm[2,2]))
#(recall= (cm_glm[2,2])/(cm_glm[2,2]+cm_glm[2,1]))
#(precision=(cm_glm[2,2])/(cm_glm[2,2]+cm_glm[1,2]))
```

Therefore the test accuracy of the logistic model is $Accuracy = \frac{TP+TN}{TP+TN+FP+FN} = \frac{224+72}{378} = 0.7848325$ i.e 78.30%


$Precision = \frac{TP}{TP+TN} = \frac{72}{72+68}$ $Recall = \frac{TP}{TP+FN} = \frac{72}{72+14}$

```
##### ROC curve #####
```

```
rocplot <- function(pred, truth) {
  predob <- prediction(pred, truth)
  perf <- performance(predob, "tpr", "fpr")
  plot(perf)
}
```

```
# Make predictions on the testing set
#predictions.fselected <- predict(model.fselct, newdata = test_data, type = "response")
predictions.fselect= predict(model, newdata = test_data, type = "response")
# Extract the predicted probabilities
fitted <- as.numeric(predictions.fselect)
```

```
# Display the ROC plot
rocplot(fitted, test_data$TravelInsurance)
```



Logestic regression:

Feature selection

-Here Age, AnnualIncome, FamilyMembers, FrequentFlyerYes, EverTravelledAbroadYes are statically significant in determining weather the customer will purchase a travel insurance or not.

```
set.seed(100)
# Step 2: Train the logistic regression model
model.fselect= glm(TravelInsurance ~ Age + AnnualIncome + FamilyMembers + FrequentFlyer + EverTravelledAbroadYes, data = train_data, family = binomial)
summary(model.fselect)

##
## Call:
## glm(formula = TravelInsurance ~ Age + AnnualIncome + FamilyMembers + FrequentFlyer + EverTravelledAbroadYes, family = binomial, data = train_data)
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)  -4.890e+00  6.854e-01  -7.135 9.69e-13 ***
## Age           5.764e-02  2.074e-02   2.779 0.005455 **
## AnnualIncome  1.542e-06  1.893e-07   8.143 3.86e-16 ***
## FamilyMembers 1.401e-01  3.765e-02   3.720 0.000199 ***
## FrequentFlyerYes 4.606e-01  1.538e-01   2.995 0.002744 **
## EverTravelledAbroadYes 1.540e+00  1.712e-01   8.995 < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##    Null deviance: 1955.0  on 1508  degrees of freedom
## Residual deviance: 1607.4  on 1503  degrees of freedom
## AIC: 1619.4
##
## Number of Fisher Scoring iterations: 4

set.seed(100)
# Step 3: Make predictions on the testing set
predictions.fselect= predict(model.fselect, newdata = test_data, type = "response")
predictions.fselect[1:10] #let's look at the first 10 predictions by the logistic model on the test data

##           8           9           10           27           39           40           41           45
## 0.8019987 0.8581790 0.2944717 0.2149270 0.7688422 0.4000402 0.7737694 0.1266537
##           46           49
## 0.7187463 0.7158887


set.seed(100)
# let's give the predicted model a good name of labels
# Convert predicted probabilities to binary predictions (0 or 1)
predicted_labels.fselect= ifelse(predictions.fselect > 0.5, 1, 0)
predicted_labels.fselect[1:10] # looking at the predictive level of first 10 observation by logistic model

##  8  9 10 27 39 40 41 45 46 49
##  1  1 0 0 1 0 1 0 1 1

#####
pred=as.factor(predicted_labels.fselect)
```

```
cm_glm.fselect = confusionMatrix(pred,test_data$TravelInsurance,positive = "1")
draw_confusion_matrix(cm_glm.fselect)
```

CONFUSION MATRIX



DETAILS

Sensitivity 0.546	Specificity 0.938	Precision 0.856	Recall 0.546	F1 0.667
Accuracy 0.78				

```
#####
set.seed(100)
cm_glm.fselect=table(predicted_labels.fselect,test_data$TravelInsurance)
(accuracy_glm.fselect=(cm_glm.fselect[1,1]+cm_glm.fselect[2,2])/(cm_glm.fselect[1,1]+cm_glm.fselect[1,2]+cm_glm.fselect[2,1]+cm_glm.fselect[2,2]))

## [1] 0.7804233

(recall_glm.fselect= (cm_glm.fselect[2,2])/(cm_glm.fselect[2,2]+cm_glm.fselect[2,1]))

## [1] 0.8556701

(precision_glm.fselect=(cm_glm.fselect[2,2])/(cm_glm.fselect[2,2]+cm_glm.fselect[1,2]))

## [1] 0.5460526


##### ROC curve #####

rocplot <- function(pred, truth) {
  predob <- prediction(pred, truth)
  perf <- performance(predob, "tpr", "fpr")
  plot(perf)
}

# Make predictions on the testing set
#predictions.fselected <- predict(model.fselct, newdata = test_data, type = "response")
predictions.fselect= predict(model.fselect, newdata = test_data, type = "response")
```

```
# Extract the predicted probabilities
fitted <- as.numeric(predictions.fselect)

# Display the ROC plot
rocplot(fitted, test_data$TravelInsurance)
```



Logistic regression

cross validation

```
set.seed(100)
logistic_regression_caret_model = train(
  form = TravelInsurance ~ Age + AnnualIncome + FamilyMembers + FrequentFlyer + EverTravelledAbroad,
  #tuneLength=10,
  data = train_data,
  trControl = trainControl(method = "cv", number = 10),
  method = "glm",
  family = "binomial"
)
summary(logistic_regression_caret_model)
```

```
##
## Call:
## NULL
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)   -4.890e+00  6.854e-01  -7.135 9.69e-13 ***
## Age             5.764e-02  2.074e-02   2.779 0.005455 **
## AnnualIncome    1.542e-06  1.893e-07   8.143 3.86e-16 ***
## FamilyMembers   1.401e-01  3.765e-02   3.720 0.000199 ***
```

```
## FrequentFlyerYes      4.606e-01  1.538e-01  2.995 0.002744 **
## EverTravelledAbroadYes 1.540e+00  1.712e-01  8.995 < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 1955.0  on 1508  degrees of freedom
## Residual deviance: 1607.4  on 1503  degrees of freedom
## AIC: 1619.4
##
## Number of Fisher Scoring iterations: 4
```

Linear Discriminant Analysis

```
set.seed(100)
#install.packages("MASS")
library(MASS)



lda.out=lda(TravelInsurance ~ Age + Employment.Type + GraduateOrNot + AnnualIncome + FamilyMembers + ChronicDiseases1 + FrequentFlyerYes + EverTravelledAbroad, data = train_data)
lda.out

## Call:
## lda(TravelInsurance ~ Age + Employment.Type + GraduateOrNot +
##      AnnualIncome + FamilyMembers + ChronicDiseases1 + FrequentFlyer +
##      EverTravelledAbroad, data = train_data)
##
## Prior probabilities of groups:
##      0      1
## 0.6494367 0.3505633
##
## Group means:
##      Age Employment.TypePrivate Sector/Self Employed GraduateOrNotYes
## 0 29.54388                                0.655102            0.8428571
## 1 29.80340                                0.805293            0.8506616
##      AnnualIncome FamilyMembers ChronicDiseases1 FrequentFlyerYes
## 0      820459.2      4.662245      0.2836735      0.1408163
## 1     1120604.9      4.918715      0.2911153      0.3327032
##      EverTravelledAbroadYes
## 0              0.06938776
## 1              0.39508507
##
## Coefficients of linear discriminants:
##
##                                LD1
## Age                        5.066645e-02
## Employment.TypePrivate Sector/Self Employed 1.521348e-01
## GraduateOrNotYes          -2.087685e-01
## AnnualIncome              1.407870e-06
## FamilyMembers             1.252036e-01
## ChronicDiseases1          2.917388e-02
## FrequentFlyerYes          4.179604e-01
## EverTravelledAbroadYes    1.717812e+00
```


Observation

-The LDA output indicates that 64.09% of the training observation corresponds to customer not taking the travel insurance and 35.90% of the training observation corresponds to the customer taking the travel insurance

```
plot(lda.out)
```



```
set.seed(100)
lda.pred <- predict(lda.out , test_data)
names(lda.pred)
```

```
## [1] "class"      "posterior" "x"
```

```
lda.pred$class[1:10] # What LDA predict for first 10 observation
```

```
## [1] 1 1 0 0 1 0 1 0 1 1
## Levels: 0 1
```


```
lda.class <- lda.pred$class
table(lda.class, test_data$TravelInsurance)
```

```
##
## lda.class   0    1
##           0 210  69
##           1  16  83
```

```
#####
```

```
cm_lda = confusionMatrix(lda.class, test_data$TravelInsurance, positive = "1")
draw_confusion_matrix(cm_lda)
```

CONFUSION MATRIX



DETAILS

Sensitivity 0.546	Specificity 0.929	Precision 0.838	Recall 0.546	F1 0.661
Accuracy 0.775				

#####

For the LDA

$$Accuracy = \frac{TP+TN}{TP+TN+FP+FN} = \frac{210+83}{378} = 0.7751 \text{ i.e } 77.51\%$$

$$Precision = \frac{TP}{TP+TN} = \frac{83}{83+16} = 0.8384 \quad Recall = \frac{TP}{TP+FN} = \frac{83}{83+69} = 0.5355$$

feature selected

```
lda.select=lda(TravelInsurance ~ Age + AnnualIncome + FamilyMembers + FrequentFlyer + EverTravelledAbroad)
lda.out
```

```
## Call:
## lda(TravelInsurance ~ Age + Employment.Type + GraduateOrNot +
##   AnnualIncome + FamilyMembers + ChronicDiseases + FrequentFlyer +
##   EverTravelledAbroad, data = train_data)
##
## Prior probabilities of groups:
##      0      1
## 0.6494367 0.3505633
##
## Group means:
##      Age Employment.Type Private Sector/Self Employed GraduateOrNotYes
## 0 29.54388                                0.655102            0.8428571
## 1 29.80340                                0.805293            0.8506616
##   AnnualIncome FamilyMembers ChronicDiseases1 FrequentFlyerYes
## 0   820459.2      4.662245      0.2836735      0.1408163
## 1  1120604.9      4.918715      0.2911153      0.3327032
##   EverTravelledAbroadYes
```

```

## 0          0.06938776
## 1          0.39508507
##
## Coefficients of linear discriminants:
##
##                               LD1
## Age                          5.066645e-02
## Employment.TypePrivate Sector/Self Employed 1.521348e-01
## GraduateOrNotYes             -2.087685e-01
## AnnualIncome                  1.407870e-06
## FamilyMembers                 1.252036e-01
## ChronicDiseases1              2.917388e-02
## FrequentFlyerYes              4.179604e-01
## EverTravelledAbroadYes        1.717812e+00
lda.pred.select <- predict(lda.select , test_data)
names(lda.pred.select)


## [1] "class"      "posterior" "x"
lda.pred.select$class[1:10] # What LDA predict for first 10 observation

## [1] 1 1 0 0 1 0 1 0 1 1
## Levels: 0 1
lda.class.select <- lda.pred.select$class
table(lda.class.select, test_data$TravelInsurance)

##
## lda.class.select    0    1
##                   0 211  70
##                   1  15  82
#####
cm_lda.select = confusionMatrix(lda.class.select, test_data$TravelInsurance, positive = "1")
draw_confusion_matrix(cm_lda.select)

```

CONFUSION MATRIX



DETAILS

Sensitivity 0.539	Specificity 0.934	Precision 0.845	Recall 0.539	F1 0.659
Accuracy 0.775				

#####

For the LDA

$$Accuracy = \frac{TP+TN}{TP+TN+FP+FN} = \frac{210+83}{378} = 0.7751 \text{ i.e } 77.51\%$$

$$Precision = \frac{TP}{TP+TN} = \frac{82}{82+70} = 0.5395 \quad Recall = \frac{TP}{TP+FN} = \frac{82}{82+15} = 0.8454$$

4) Quadratic Discriminant Analysis

```
qda.out=qda(TravelInsurance ~ Age + Employment.Type + GraduateOrNot + AnnualIncome + FamilyMembers + Ch
qda.out
```

```
## Call:
## qda(TravelInsurance ~ Age + Employment.Type + GraduateOrNot +
##   AnnualIncome + FamilyMembers + ChronicDiseases + FrequentFlyer +
##   EverTravelledAbroad, data = train_data)
##
## Prior probabilities of groups:
##      0      1
## 0.6494367 0.3505633
##
## Group means:
##      Age Employment.Type Private Sector/Self Employed GraduateOrNotYes
## 0 29.54388                                0.655102            0.8428571
## 1 29.80340                                0.805293            0.8506616
##   AnnualIncome FamilyMembers ChronicDiseases1 FrequentFlyerYes
## 0   820459.2      4.662245      0.2836735      0.1408163
## 1  1120604.9      4.918715      0.2911153      0.3327032
##   EverTravelledAbroadYes
## 0              0.06938776
```

```
## 1 0.39508507
```

Observation -The QDA output indicates that 64.09% of the training observation corresponds to customer not taking the travel insurance and 35.90% of the training observation corresponds to the customer taking the travel insurance

```
qda.pred <- predict(qda.out , test_data)
names(qda.pred)
```

```
## [1] "class"      "posterior"
```

```
qda.pred$class[1:10] # What QDA predict for first 10 observation
```

```
## [1] 1 1 0 0 1 0 1 0 1 1
```

```
## Levels: 0 1
```

```
set.seed(100)
```

```
qda.class <- qda.pred$class
```

```
table(qda.class, test_data$TravelInsurance)
```

```
##
```

```
## qda.class  0  1
```

```
##          0 204  62
```


```
##          1  22  90
```

```
#####
```

```
cm_qda = confusionMatrix(qda.class, test_data$TravelInsurance, positive = "1")
```

```
draw_confusion_matrix(cm_qda )
```

CONFUSION MATRIX



DETAILS

Sensitivity 0.592	Specificity 0.903	Precision 0.804	Recall 0.592	F1 0.682
Accuracy 0.778				

```
#####
```

```
(204+90)/378
```

```
## [1] 0.7777778
```

```
90/(90+62)
```

```
## [1] 0.5921053
```

```
90/(90+22)
```

```
## [1] 0.8035714
```

$Accuracy = \frac{TP+TN}{TP+TN+FP+FN} (204+90)/378 = 0.7778$ i.e 77.78%

$Precision = \frac{TP}{TP+FN} = \frac{90}{90+62} = 0.5921$ $Recall = \frac{TP}{TP+FN} = \frac{90}{90+22} = 0.8036$

feature selection

```
qda.select=qda(TravelInsurance ~ Age + AnnualIncome + FamilyMembers + FrequentFlyer + EverTravelledAbroad, data = train_data)
qda.select
```

```
## Call:
```

```
## qda(TravelInsurance ~ Age + AnnualIncome + FamilyMembers + FrequentFlyer + EverTravelledAbroad, data = train_data)
```

```
##
```

```
## Prior probabilities of groups:
```

```
##      0      1
```

```
## 0.6494367 0.3505633
```

```
##
```

```
## Group means:
```

```
##      Age AnnualIncome FamilyMembers FrequentFlyerYes EverTravelledAbroadYes
```

```
## 0 29.54388      820459.2      4.662245      0.1408163      0.06938776
```

```
## 1 29.80340      1120604.9      4.918715      0.3327032      0.39508507
```

```
qda.pred.select <- predict(qda.select, test_data)
```

```
names(qda.pred.select)
```

```
## [1] "class"      "posterior"
```

```
qda.pred.select$class[1:10] # What QDA predict for first 10 observation
```

```
## [1] 1 1 0 0 1 1 1 0 1 1
```

```
## Levels: 0 1
```

```
set.seed(100)
```

```
qda.class.select <- qda.pred.select$class
```

```
table(qda.class.select, test_data$TravelInsurance)
```

```
##
```

```
## qda.class.select      0      1
```

```
##      0 203  63
```

```
##      1  23  89
```

```
#####
```

```
cm_qda.select = confusionMatrix(qda.class.select, test_data$TravelInsurance, positive = "1")
draw_confusion_matrix(cm_qda.select)
```

CONFUSION MATRIX

		Actual	
		0	1
Predicted	0	203	63
	1	23	89

DETAILS

Sensitivity 0.586	Specificity 0.898	Precision 0.795	Recall 0.586	F1 0.674
Accuracy 0.772				

```
#####
```

```
(203+89)/378
```

```
## [1] 0.7724868
```

```
89/(89+63)
```

```
## [1] 0.5855263
```

```
89/(89+23)
```

```
## [1] 0.7946429
```

```
Accuracy =  $\frac{TP+TN}{TP+TN+FP+FN}$  (203+89)/378 =0.7725 i.e 77.25%
```

```
Precision =  $\frac{TP}{TP+FN}$  =  $\frac{89}{89+63}$  =0.5855 Recall =  $\frac{TP}{TP+FN}$  =  $\frac{89}{89+23}$  = 0.7946
```

```
KNN ##### 1) KNN- classifier
```

```
set.seed(100)
```

```
# categorical variable as factor
```

```
Insurance_data$ChronicDiseases=as.factor(Insurance_data$ChronicDiseases)
```

```
Insurance_data$Employment.Type= as.factor(Insurance_data$Employment.Type)
```

```
Insurance_data$GraduateOrNot= as.factor(Insurance_data$GraduateOrNot)
```

```
Insurance_data$FrequentFlyer= as.factor(Insurance_data$FrequentFlyer)
```

```
Insurance_data$EverTravelledAbroad= as.factor(Insurance_data$EverTravelledAbroad)
```

```
Insurance_data$TravelInsurance= as.factor(Insurance_data$TravelInsurance)
```

```
# converting all my dataset to numeric for the model setting
```

```
Insurance_data_num <- as.data.frame(lapply(Insurance_data[,1:8], as.numeric))
```

```

set.seed(100)
knn_fit = train(
  TravelInsurance ~ .,
  data = train_data,
  method = "knn",
  tuneLength=10,
  trControl = trainControl(method = "cv", number = 10),
  preProcess = c("center", "scale")
)

knn_fit

## k-Nearest Neighbors
##
## 1509 samples
##      8 predictor
##      2 classes: '0', '1'
##
## Pre-processing: centered (8), scaled (8)
## Resampling: Cross-Validated (10 fold)
## Summary of sample sizes: 1358, 1358, 1358, 1358, 1358, 1358, ...
## Resampling results across tuning parameters:
##
##  k  Accuracy  Kappa
##   5  0.7621148  0.4528058
##   7  0.7720618  0.4654331
##   9  0.7713687  0.4556991
##  11  0.7753687  0.4626138
##  13  0.7786799  0.4676690
##  15  0.7846402  0.4800932
##  17  0.7866402  0.4842670
##  19  0.7839912  0.4807701
##  21  0.7813377  0.4725415
##  23  0.7853157  0.4792685
##
## Accuracy was used to select the optimal model using the largest value.
## The final value used for the model was k = 17.

KNN_pred=predict(knn_fit , test_data)
cm=table(KNN_pred,test_data$TravelInsurance)
(212+94)/378

## [1] 0.8095238
94/(94+58)

## [1] 0.6184211
94/(94+14)

## [1] 0.8703704
#####
cm_KNN = confusionMatrix(KNN_pred,test_data$TravelInsurance,positive = "1")
draw_confusion_matrix(cm_KNN)

```


CONFUSION MATRIX

		Actual	
		0	1
Predicted	0	212	59
	1	14	93

DETAILS

Sensitivity 0.612	Specificity 0.938	Precision 0.869	Recall 0.612	F1 0.718
Accuracy 0.807				

#####

$Accuracy = \frac{TP+TN}{TP+TN+FP+FN} = \frac{169+39}{378} = 0.8095$ i.e 80.95%

$Precision = \frac{TP}{TP+TN} = \frac{94}{94+58} = 0.6184$ $Recall = \frac{TP}{TP+FN} = \frac{94}{94+14} = 0.8704$

using the cross validation for finding the best number of class, k=17

DT and Pruned DT


```
tree.d=tree(TravelInsurance~., data=Insurance_data,split="gini",subset= sample_index)
summary(tree.d)
```

```
##
## Classification tree:
## tree(formula = TravelInsurance ~ ., data = Insurance_data, subset = sample_index,
##       split = "gini")
## Number of terminal nodes: 154
## Residual mean deviance: 0.763 = 1034 / 1355
## Misclassification error rate: 0.1637 = 247 / 1509
```

- The Insurance tree has 154 terminal nodes or leaves, which are the endpoints where the classification decisions are made. It is very crowded tree
- The residual mean deviance is 0.763. A lower deviance indicates a better fit of the model to the data.
- The misclassification error rate for this tree is 0.164, which is calculated as 247 misclassified cases out of a total of 159 cases in train data.

(2) Create a plot of the tree. Pick one of the terminal nodes, and interpret the information displayed.

```
plot(tree.d)
```



```
#text(tree.d)  
#tree.d
```

- (3) Predict the labels on the test data, and produce a confusion matrix comparing the test labels to the predicted test labels. What is the test error rate?

```
set.seed(100)  
pred.d=predict(tree.d,test_data,type="class")  
  
DT.cm=confusionMatrix(pred.d, test_data$TravelInsurance,positive = "1")  
draw_confusion_matrix(DT.cm)
```

CONFUSION MATRIX

		Actual	
		0	1
Predicted	0	209	51
	1	17	101

DETAILS


Sensitivity 0.664	Specificity 0.925	Precision 0.856	Recall 0.664	F1 0.748
Accuracy 0.82				

```
table(pred.d, test_data$TravelInsurance)
```

```
##
## pred.d  0   1
##        0 209  51
##        1  17 101
```

- (4) Apply the `cv.tree()` function to the training set in order to determine the optimal tree size. Produce a plot with tree size on the x-axis and cross-validated classification error rate on the y-axis. Which tree size corresponds to the lowest cross-validated classification error rate?

```
#pruning
cv.d=cv.tree(tree.d)
plot(cv.d$size, cv.d$dev, type="b")
```



```
cv.d$dev
```

```
## [1] 1363.359 1363.359 1363.359 1363.359 1363.359 1363.359 1363.359 1363.359 1363.359
## [9] 1363.359 1363.359 1363.359 1363.359 1363.359 1363.359 1363.359 1363.359 1363.359
## [17] 1363.359 1363.359 1363.359 1363.359 1363.359 1363.359 1363.359 1363.359 1363.359
## [25] 1363.359 1363.359 1363.359 1363.359 1363.359 1363.359 1363.359 1363.359 1363.359
## [33] 1363.359 1363.359 1363.359 1363.359 1363.359 1363.359 1363.359 1363.359 1363.359
## [41] 1363.359 1363.359 1363.359 1363.359 1363.359 1363.359 1363.359 1363.359 1363.359
## [49] 1363.359 1363.359 1363.359 1363.359 1363.359 1363.359 1363.359 1363.359 1363.359
## [57] 1363.359 1363.359 1363.359 1363.359 1363.359 1363.359 1363.359 1363.359 1363.359
## [65] 1363.359 1363.359 1363.359 1363.359 1363.359 1363.359 1363.359 1363.359 1363.359
## [73] 1363.359 1363.359 1363.359 1363.359 1363.359 1363.359 1363.359 1363.359 1363.359
## [81] 1363.359 1363.359 1363.359 1363.359 1363.359 1363.359 1363.359 1363.359 1363.359
## [89] 1368.298 1402.694 1572.127 1956.913
```


Because deviance error is constant after tree size = 5, I chose tree size = 5.

- (5) Produce a pruned tree corresponding to the optimal tree size obtained using cross-validation. If cross-validation does not lead to selection of a pruned tree, then create a pruned tree with five terminal nodes.

```
prune.d=prune.tree(tree.d,best=5)
summary(prune.d)
```

```
##
## Classification tree:
## snip.tree(tree = tree.d, nodes = c(7L, 6L, 11L, 10L, 4L))
## Variables actually used in tree construction:
## [1] "AnnualIncome" "Age" "FamilyMembers"
## Number of terminal nodes: 5
## Residual mean deviance: 0.9015 = 1356 / 1504
## Misclassification error rate: 0.1723 = 260 / 1509
```

```
plot(prune.d)
text(prune.d)
```



for Best tree size = 5:

Residual mean deviance: $0.901 = 1360 / 1500$ Misclassification error rate: $0.172 = 260 / 1509$

Both Residual mean deviance and Misclassification error rate are greater for best tree size = 5

(6) Compare the training and test error rates between the pruned and unpruned trees. Which is higher?

```

set.seed(100)
pred.prune=predict(prune.d,test_data,type="class")

prune_DT.cm=confusionMatrix(pred.prune, test_data$TravelInsurance,positive = "1")
draw_confusion_matrix(prune_DT.cm)
  
```

CONFUSION MATRIX

		Actual	
		0	1
Predicted	0	219	53
	1	7	99

DETAILS

Sensitivity	Specificity	Precision	Recall	F1
0.651	0.969	0.934	0.651	0.767
Accuracy 0.841				

```
table(pred.d,test_data$TravelInsurance)
```

```
##
## pred.d    0    1
##          0 209  51
##          1   7 101
```

```
(test.error.DT=(42+26)/(270))
```

```
## [1] 0.2518519
```

```
(test.error.prune=(31+32)/270)
```

```
## [1] 0.2333333
```

For DT: Residual mean deviance: 0.6359 = 455.3 / 716 Misclassification error rate: 0.1525 = 122 / 800
 $\text{test.error.DT} = 0.251 = (42+26)/(270)$

For Pruned DT: Residual mean deviance: 1.088 = 863.5 / 794 Misclassification error rate: 0.2788 = 223 / 800
 $\text{test.error.prune} = 0.233 = (31+32)/270$

The test error for pruned DT is less than the test error for unpruned dt which is predictable.

Random Forest

```
set.seed(100)
bag.Insurance_data=randomForest(TravelInsurance~., data=Insurance_data, subset= sample_index, mtry=8, i
bag.Insurance_data # lets take a look at the output
```

```
##
## Call:
```

```
## randomForest(formula = TravelInsurance ~ ., data = Insurance_data, mtry = 8, importance = TRUE)
##           Type of random forest: classification
##           Number of trees: 500
## No. of variables tried at each split: 8
##
##           OOB estimate of  error rate: 21.6%
## Confusion matrix:
##      0   1 class.error
## 0 860 120   0.122449
## 1 206 323   0.389414
```


```
yhat.bag=predict(bag.Insurance_data, newdata = test_data)
```

```
table(yhat.bag, test_data$TravelInsurance)
```

```
##
## yhat.bag   0   1
##           0 207  50
##           1  19 102
```

```
bag.cm=confusionMatrix(yhat.bag, test_data$TravelInsurance,positive = "1")
draw_confusion_matrix(bag.cm)
```

CONFUSION MATRIX



DETAILS

Sensitivity 0.671	Specificity 0.916	Precision 0.843	Recall 0.671	F1 0.747
Accuracy 0.817				

```
(mtry=round(sqrt(8),0))
```

```
## [1] 3
```

```
# best mtry =3
```

```
set.seed(100)
```

```
(rf.fit3=randomForest(TravelInsurance~., data=Insurance_data, subset= sample_index, ntree=1000, mtry=3,
```

```
##
```

```
## Call:
```

```
## randomForest(formula = TravelInsurance ~ ., data = Insurance_data, ntree = 1000, mtry = 3, imp
```

```
## Type of random forest: classification
```

```
## Number of trees: 1000
```

```
## No. of variables tried at each split: 3
```

```
##
```

```
## OOB estimate of error rate: 18.42%
```

```
## Confusion matrix:
```

```
## 0 1 class.error
```

```
## 0 930 50 0.05102041
```

```
## 1 228 301 0.43100189
```

```
importance(bag.Insurance_data)
```

	0	1	MeanDecreaseAccuracy	MeanDecreaseGini
## Age	31.808691	43.947378	51.1346503	83.184216
## Employment.Type	17.209817	3.879430	18.3146677	12.742815
## GraduateOrNot	12.091890	9.076818	15.8103116	8.623758
## AnnualIncome	75.650550	112.331823	124.9789761	286.112323
## FamilyMembers	22.568557	53.404849	47.3448836	132.096367
## ChronicDiseases	-4.193267	6.575596	0.4094072	28.998157
## FrequentFlyer	1.646629	3.999756	3.9853358	19.094321
## EverTravelledAbroad	7.940023	1.606792	7.7624097	11.239449

```
importance(rf.fit3)
```

	0	1	MeanDecreaseAccuracy	MeanDecreaseGini
## Age	47.862392	57.836976	73.699745	79.49270
## Employment.Type	19.183589	12.656292	24.180313	11.92546
## GraduateOrNot	8.011737	13.316647	14.974379	8.75125
## AnnualIncome	40.587780	130.352941	104.434678	198.70659
## FamilyMembers	38.907573	64.885550	66.955196	82.97734
## ChronicDiseases	-4.406492	6.706806	0.222829	16.95500
## FrequentFlyer	9.157400	24.777226	24.339995	18.55546
## EverTravelledAbroad	9.242002	33.824674	35.117362	49.18150

AnnualIncome, Age, FamilyMembers, Employment.Type, EverTravelledAbroadYes are important variables in Bagging and RF.

```
set.seed(100)
```

```
yhat.RF=predict(rf.fit3, newdata = test_data)
```

```
table(yhat.RF, test_data$TravelInsurance)
```

```
##
```

```
## yhat.RF 0 1
```

```
## 0 219 46
```

```
## 1 7 106
```

```
RF.cm=confusionMatrix(yhat.RF, test_data$TravelInsurance,positive = "1")
```

```
draw_confusion_matrix(RF.cm)
```


CONFUSION MATRIX

		Actual	
		0	1
Predicted	0	219	46
	1	7	106

DETAILS

Sensitivity 0.697	Specificity 0.969	Precision 0.938	Recall 0.697	F1 0.8
Accuracy 0.86				

```
TravelInsuranceTest$TravelInsurance=rep(0,100)
TravelInsuranceTest$TravelInsurance=as.factor(TravelInsuranceTest$TravelInsurance)
yhat.RF=predict(rf.fit3, newdata = TravelInsuranceTest)
table(yhat.RF)
```

```
## yhat.RF
## 0 1
## 78 22
```

```
TravelInsuranceTest$TravelInsurance=yhat.RF
#write.csv(TravelInsuranceTest,"TravelInsuranceTest_Labeled.csv")
#read.csv("TravelInsuranceTest_Labeled.csv",header = T)
```

XGBoost

hyperparameter tuning

```
set.seed(100)

grid_gbm = expand.grid(nrounds = c(1,10),
  max_depth = c(1,4),
  eta = c(.1,.4),
  gamma = 0,
  colsample_bytree = .7,
  min_child_weight = 1,
  subsample = c(.8,1))
```


CONFUSION MATRIX

		Actual	
		0	1
Predicted	0	220	54
	1	6	98

DETAILS

Sensitivity 0.645	Specificity 0.973	Precision 0.942	Recall 0.645	F1 0.766
Accuracy 0.841				

SVM

- (d) Tune the linear SVM with various values of cost. Report the cross-validation errors associated with different values of this parameter. Select an optimal cost. Compute the training and test error rates using this new cost value. Comment on your findings.

```
set.seed(100)
tune.out=tune(svm,TravelInsurance~.,data = train_data,kernel="linear",scale=T,
              ranges=list(cost=c(0.001, 0.01, 0.1,0.5, 1,2,5,10,100)))
summary(tune.out)
```

```
##
## Parameter tuning of 'svm':
##
## - sampling method: 10-fold cross validation
##
## - best parameters:
##   cost
##   0.01
##
## - best performance: 0.2366269
##
## - Detailed performance results:
##   cost      error dispersion
## 1 1e-03 0.3505784 0.03920048
## 2 1e-02 0.2366269 0.04428797
## 3 1e-01 0.2571567 0.04059054
## 4 5e-01 0.2571567 0.04059054
```

```

## 5 1e+00 0.2571567 0.04059054
## 6 2e+00 0.2571567 0.04059054
## 7 5e+00 0.2571567 0.04059054
## 8 1e+01 0.2571567 0.04059054
## 9 1e+02 0.2571567 0.04059054

tune.out$best.parameters

## cost
## 2 0.01

tune.out$best.performance

## [1] 0.2366269

set.seed(100)
best.fit = svm(TravelInsurance~.,data = train_data, kernel = "linear", cost = 0.01, scale = TRUE)

# best fit training error rate
pred_train=predict(best.fit , train_data)
table(pred_train, train_data$TravelInsurance)

##
## pred_train  0  1
##           0 919 301
##           1  61 228

# best performance error : 0.2366
best.train.err=(301+61)/1506

# best fittest error rate
pred_test=predict(best.fit , test_data)
table(pred_test, test_data$TravelInsurance)

##
## pred_test  0  1
##           0 215  71
##           1  11  81

(best.test.err = (71+11)/378)

## [1] 0.2169312

#0.2169

#####
cm_svm = confusionMatrix(pred_test,test_data$TravelInsurance,positive = "1")
draw_confusion_matrix(cm_svm)

```

CONFUSION MATRIX

		Actual	
		0	1
Predicted	0	215	71
	1	11	81

DETAILS

Sensitivity 0.533	Specificity 0.951	Precision 0.88	Recall 0.533	F1 0.664
Accuracy 0.783				

#####

finding the best cost using crossvalidation and fit the model

we tuned the svm model with 10-fold cross validation, the best parameter for cost =0.01, and best performance error = 0.2366

The error rate in the training data is 0.24045, and in the test data is 0.2169.

(e) Now repeat (d), with radial basis kernels, with different values of gamma and cost. Comment on your results. Which approach seems to give the better results on this data?

```
set.seed(100)
# finding best values of gamma and cost
tune.out=tune(svm,TravelInsurance~.,data = train_data, kernel="radial",ranges=list(cost=c(0.001, 0.01, 0.1),gamma=c(0.001, 0.01, 0.1)))
summary(tune.out)
```

```
##
## Parameter tuning of 'svm':
##
## - sampling method: 10-fold cross validation
##
## - best parameters:
##   cost gamma
##   1 0.5
##
## - best performance: 0.1856203
##
## - Detailed performance results:
##   cost gamma   error dispersion
## 1 1e-03 0.1 0.3505784 0.03920048
```

```
## 2 1e-02 0.1 0.3505784 0.03920048
## 3 1e-01 0.1 0.2200706 0.04793568
## 4 1e+00 0.1 0.1915717 0.04585993
## 5 5e+00 0.1 0.1862781 0.05114341
## 6 1e+01 0.1 0.1862781 0.04748726
## 7 1e-03 0.5 0.3505784 0.03920048
## 8 1e-02 0.5 0.3505784 0.03920048
## 9 1e-01 0.5 0.2081413 0.04608717
## 10 1e+00 0.5 0.1856203 0.05170815
## 11 5e+00 0.5 0.1915585 0.03756448
## 12 1e+01 0.5 0.1895585 0.02890363
## 13 1e-03 1.0 0.3505784 0.03920048
## 14 1e-02 1.0 0.3505784 0.03920048
## 15 1e-01 1.0 0.2445784 0.03919909
## 16 1e+00 1.0 0.1915806 0.04831654
## 17 5e+00 1.0 0.2101060 0.03719872
## 18 1e+01 1.0 0.2273377 0.03336786
## 19 1e-03 2.0 0.3505784 0.03920048
## 20 1e-02 2.0 0.3505784 0.03920048
## 21 1e-01 2.0 0.3293819 0.04404555
## 22 1e+00 2.0 0.2127903 0.04797584
## 23 5e+00 2.0 0.2432539 0.04318160
## 24 1e+01 2.0 0.2512009 0.03840622
## 25 1e-03 3.0 0.3505784 0.03920048
## 26 1e-02 3.0 0.3505784 0.03920048
## 27 1e-01 3.0 0.3419647 0.04064535
## 28 1e+00 3.0 0.2240442 0.04551785
## 29 5e+00 3.0 0.2452362 0.03782538
## 30 1e+01 3.0 0.2465651 0.03935492
## 31 1e-03 4.0 0.3505784 0.03920048
## 32 1e-02 4.0 0.3505784 0.03920048
## 33 1e-01 4.0 0.3479294 0.03715165
## 34 1e+00 4.0 0.2379603 0.04707405
## 35 5e+00 4.0 0.2498852 0.04174338
## 36 1e+01 4.0 0.2498852 0.04266707
```

```
tune.out$best.parameters
```

```
##      cost gamma
## 10      1    0.5
```

```
tune.out$best.performance
```

```
## [1] 0.1856203
```

```
radial.svmfit = svm(TravelInsurance~.,data = train_data, kernel = "radial",gamma=0.1, cost = 10, decision
```

```
# training error rate
```

```
radial.pred_train =predict(radial.svmfit , train_data)
table(radial.pred_train , train_data$TravelInsurance)
```

```
##
## radial.pred_train  0  1
##                   0 943 231
##                   1  37 298
```

```

(radial.train.err = (231+37)/1509)


## [1] 0.1776011
# test error rate
radial.pred_test=predict(radial.svmfit, test_data)
table(radial.pred_test, test_data$TravelInsurance)

##
## radial.pred_test  0   1
##                  0 216  56
##                  1  10  96
(radial.test.err = (56+10)/378)

## [1] 0.1746032
cm_svm_radial = confusionMatrix(radial.pred_test,test_data$TravelInsurance,positive = "1")
draw_confusion_matrix(cm_svm_radial)

```

CONFUSION MATRIX



DETAILS

Sensitivity 0.632	Specificity 0.956	Precision 0.906	Recall 0.632	F1 0.744
Accuracy 0.825				

permormance error: 0.1829 cost=10 gamma =0.1 The training error for the radial kernel (0.1776) is lower than that of the linear kernel (0.24045). However, the test error for the radial kernel (0.1746) less than the linear kernel (0.2169). Therefore, based on these results, it appears that the redial kernel is more effective for our dataset.

- (f) Now repeat again, with polynomial basis kernels, with different values of degree and cost. Comment on your results. Which approach (kernel) seems to give the best results on this data?

```

set.seed(100)
# finding best values of gamma and cost
tune.out=tune(svm,TravelInsurance~.,data = train_data, kernel="polynomial",ranges=list(cost=c(0.001, 0.

```

```
summary(tune.out)
```

```
##
## Parameter tuning of 'svm':
##
## - sampling method: 10-fold cross validation
##
## - best parameters:
##   cost degree
##     5       3
##
## - best performance: 0.1856159
##
## - Detailed performance results:
##   cost degree      error dispersion
## 1  1e-03    0.1 0.3505784 0.03920048
## 2  1e-02    0.1 0.3505784 0.03920048
## 3  1e-01    0.1 0.3505784 0.03920048
## 4  1e+00    0.1 0.3505784 0.03920048
## 5  5e+00    0.1 0.3505784 0.03920048
## 6  1e+01    0.1 0.3505784 0.03920048
## 7  1e+02    0.1 0.3505784 0.03920048
## 8  1e-03    0.5 0.3505784 0.03920048
## 9  1e-02    0.5 0.3505784 0.03920048
## 10 1e-01    0.5 0.3505784 0.03920048
## 11 1e+00    0.5 0.3505784 0.03920048
## 12 5e+00    0.5 0.3505784 0.03920048
## 13 1e+01    0.5 0.3505784 0.03920048
## 14 1e+02    0.5 0.3505784 0.03920048
## 15 1e-03    1.0 0.3505784 0.03920048
## 16 1e-02    1.0 0.3505784 0.03920048
## 17 1e-01    1.0 0.2445784 0.04605799
## 18 1e+00    1.0 0.2571567 0.04059054
## 19 5e+00    1.0 0.2571567 0.04059054
## 20 1e+01    1.0 0.2571567 0.04059054
## 21 1e+02    1.0 0.2571567 0.04059054
## 22 1e-03    2.0 0.3505784 0.03920048
## 23 1e-02    2.0 0.3505784 0.03920048
## 24 1e-01    2.0 0.2233731 0.04642170
## 25 1e+00    2.0 0.2227196 0.04492097
## 26 5e+00    2.0 0.2154305 0.04415458
## 27 1e+01    2.0 0.2134481 0.04459755
## 28 1e+02    2.0 0.2127859 0.04445315
## 29 1e-03    3.0 0.3505784 0.03920048
## 30 1e-02    3.0 0.3505784 0.03920048
## 31 1e-01    3.0 0.2339823 0.05184210
## 32 1e+00    3.0 0.1922428 0.04631032
## 33 5e+00    3.0 0.1856159 0.04905982
## 34 1e+01    3.0 0.1882649 0.04578042
## 35 1e+02    3.0 0.1909007 0.04263200
## 36 1e-03    4.0 0.3505784 0.03920048
## 37 1e-02    4.0 0.3505784 0.03920048
## 38 1e-01    4.0 0.2452362 0.04983317
## 39 1e+00    4.0 0.2008477 0.04549404
```



```
## 40 5e+00    4.0 0.1875938 0.04590359
## 41 1e+01    4.0 0.1862781 0.04758977
## 42 1e+02    4.0 0.1929007 0.05252725
```

```
tune.out$best.parameters
```

```
##      cost degree
## 33      5      3
```

```
tune.out$best.performance
```

```
## [1] 0.1856159
```

```
poly.svmfit = svm(TravelInsurance~.,data = train_data, kernel = "radial",degree=3, cost = 5, decision.v
# traning error rate
poly.pred_train =predict(poly.svmfit , train_data)
table(poly.pred_train , train_data$TravelInsurance)
```

```
##
## poly.pred_train    0    1
##                   0 942 233
##                   1  38 296
```

```
(poly.train.err = (233+38)/1509)
```

```
## [1] 0.1795891
```

```
# test error rate
poly.pred_test=predict(poly.svmfit, test_data)
table(poly.pred_test, test_data$TravelInsurance)
```


```
##
## poly.pred_test    0    1
##                   0 214  56
##                   1  12  96
```

```
(poly.test.err = (56+12)/378)
```

```
## [1] 0.1798942
```

```
#####
cm_svm_poly = confusionMatrix(poly.pred_test,test_data$TravelInsurance,positive = "1")
draw_confusion_matrix(cm_svm_poly)
```

CONFUSION MATRIX



DETAILS

Sensitivity 0.632	Specificity 0.947	Precision 0.889	Recall 0.632	F1 0.738
Accuracy 0.82				

```
#####
```

```
c(best.train.err,best.test.err)
```

```
## [1] 0.2403718 0.2169312
```

```
c(radial.train.err,radial.test.err)
```

```
## [1] 0.1776011 0.1746032
```

```
c(poly.train.err,poly.test.err)
```

```
## [1] 0.1795891 0.1798942
```

cost = 5 degree = 3 performance error = 0.1856 comparing the training and test error for linear, radial, and polynomial kernels we can see that radial kernel has the best performance.

Neural Network

```
standardize=function(x) {(x-min(x))/(max(x)-min(x))}
std.data=Insurance_data
std.data$AnnualIncome=standardize(std.data$AnnualIncome)
std.data$Age=standardize(std.data$Age)
std.data$FamilyMembers=standardize(std.data$FamilyMembers)
set.seed(100)
ind=sample(1:nrow(std.data), 0.8*nrow(std.data))
train=std.data[ind,]
test=std.data[-ind,]
```

```
set.seed(100)
fit=nnet(TravelInsurance~., data=train,decay=0.1, size=10, liout=FALSE)
```

```
## # weights: 101
## initial value 1122.179933
## iter 10 value 816.650248
## iter 20 value 763.223410
## iter 30 value 739.491369
## iter 40 value 725.620421
## iter 50 value 720.433385
## iter 60 value 717.795530
## iter 70 value 716.821754
## iter 80 value 716.263152
## iter 90 value 715.479459
## iter 100 value 714.797982
## final value 714.797982
## stopped after 100 iterations
```

(b) Compare the classification performance of your model with that of linear logistic regression.

```
set.seed(100)
NN_probs=predict(fit, test)
NN_pred <- rep("No",378)
NN_pred[NN_probs > 0.5] = "Yes"

# The confusion matrix
(cm <- table( NN_pred,test_data$TravelInsurance))
```

```
##
## NN_pred  0  1
##      No 215 55
##      Yes 11 97
```

```
#drawing confusion matrix
#####
NN_predicted_labels= ifelse(NN_probs > 0.5, 1, 0)

pred=as.factor(NN_predicted_labels)
cm_NN = confusionMatrix(pred,test$TravelInsurance,positive= "1")

draw_confusion_matrix(cm_NN)
```

CONFUSION MATRIX

		Actual	
		0	1
Predicted	0	215	55
	1	11	97

DETAILS

Sensitivity	Specificity	Precision	Recall	F1
0.638	0.951	0.898	0.638	0.746
Accuracy				
0.825				

```
#####
set.seed(100)
mygrid=expand.grid(.decay=c(0.05,0.1),.size=c(3,4,5,6,7,8,9,10,12))
nnetfit=train(TravelInsurance~., data=train, method= "nnet", mmaxit=1000,tuneGrid= mygrid,trace=F)
nnetfit
```

```
## Neural Network
##
## 1509 samples
## 8 predictor
## 2 classes: '0', '1'
##
## No pre-processing
## Resampling: Bootstrapped (25 reps)
## Summary of sample sizes: 1509, 1509, 1509, 1509, 1509, 1509, ...
## Resampling results across tuning parameters:
##
##  decay  size  Accuracy  Kappa
##  0.05    3    0.7774416  0.4674308
##  0.05    4    0.7826637  0.4820034
##  0.05    5    0.7808932  0.4832054
##  0.05    6    0.7785403  0.4767325
##  0.05    7    0.7784117  0.4786143
##  0.05    8    0.7799334  0.4821469
##  0.05    9    0.7774541  0.4777516
##  0.05   10    0.7810647  0.4874769
##  0.05   12    0.7740846  0.4732448
##  0.10    3    0.7764037  0.4651277
```

```
## 0.10 4 0.7822076 0.4807091
## 0.10 5 0.7839584 0.4872408
## 0.10 6 0.7840119 0.4887891
## 0.10 7 0.7827955 0.4867654
## 0.10 8 0.7794185 0.4797147
## 0.10 9 0.7791144 0.4803317
## 0.10 10 0.7794541 0.4804515
## 0.10 12 0.7797722 0.4828214
##
## Accuracy was used to select the optimal model using the largest value.
## The final values used for the model were size = 6 and decay = 0.1.

choose decay=0.1 and size=6 accuracy=0.7840
```

```
set.seed(100)
fit=nnet(TravelInsurance~., data=train,decay=0.1, size=6, liout=FALSE)
```

```
## # weights: 61
## initial value 982.764547
## iter 10 value 810.172714
## iter 20 value 758.474694
## iter 30 value 743.609932
## iter 40 value 734.590877
## iter 50 value 731.003901
## iter 60 value 728.840800
## iter 70 value 727.548934
## iter 80 value 727.253517
## iter 90 value 727.132288
## iter 100 value 727.098660
## final value 727.098660
## stopped after 100 iterations
```

(b) Compare the classification performance of your model with that of linear logistic regression.

```
set.seed(100)
NN_probs=predict(fit, test)
NN_pred <- rep("No",378)
NN_pred[NN_probs > 0.5] = "Yes"

# The confusion matrix
(cm <- table( NN_pred,test$TravelInsurance))
```

```
##
## NN_pred 0 1
## No 214 58
## Yes 12 94
```


```
#drawing confusion matrix
#####
NN_predicted_labels= ifelse(NN_probs > 0.5, 1, 0)

pred=as.factor(NN_predicted_labels)

cm_NN = confusionMatrix(pred,test$TravelInsurance,positive= "1")

draw_confusion_matrix(cm_NN)
```

CONFUSION MATRIX



DETAILS

Sensitivity 0.618	Specificity 0.947	Precision 0.887	Recall 0.618	F1 0.729
Accuracy 0.815				

#####